

# Haoran Zhang

✉ zhhaoran@umich.edu | 💬 zhhr0321

## Education

<b>University of Michigan, Ann Arbor</b> B.Eng in Computer Science	August 2024 – May 2026 Michigan, United States
• major GPA 3.98/4.0 • A+/A Courses: Intro to Operating System, Compiler, Computer Networks, Cryptography, etc.	
<b>Shanghai Jiao Tong University</b> B.Eng in Mech Engineering	September 2022 – August 2026 Shanghai, China
• A+/A Courses: Probabilistic Methods in Engineering, Honors Calculus IV, Honors Calculus III, etc.	

## Research Experience

<b>Agentic-ds-ops</b> Distributed Systems Project	MAY. 2025 – PRESENT Order Lab, Ann Arbor, MI
• Built an <b>agent-based autonomous mitigation system</b> for <b>distributed failures</b> like overload, network faults on <b>ZooKeeper</b> clusters. Experimented with a custom overload benchmark in Go.	
• Integrated <b>Prometheus</b> metrics with sliding window to dynamically detect metrics trend and <b>JMX-exporter</b> to collect logs. Combined for failure detection.	
• Designed <b>risk-aware mitigation framework</b> where agent selects from pre-defined actions via <b>Haproxy</b> , <b>Resilience4j</b> based on symptom severity.	
• Applied pre- and post-evaluator to quantitatively <b>predict</b> and <b>verify</b> outcomes against <b>SLOs/throughput</b> before concluding mitigation success.	
<b>CUDA Graphs for Reducing Kernel Launch Overhead</b> MLSys Research Project	SEP. 2025 – PRESENT Ann Arbor, MI
• Built a hybrid <b>runtime proxy</b> reducing <b>kernel launch overhead</b> and cut down on <b>tail latency</b> in Large recommendation system, MoE inference.	
• Applied <b>CUDA Graphs</b> for stable, high-arithmetic compute (MLP blocks/attention) and a <b>persistent kernel</b> consuming device-queue tasks for irregular micro-ops (pack/scatter/routing).	
• Used <b>bucketing</b> and <b>static pools</b> to capture several robust graphs and achieved higher hit rate of <b>graph replay</b> .	
<b>LLM-Powered Robotic Manipulation System</b> Robotics Research	MAR. 2024 – AUG. 2024 Shanghai Jiao Tong University
• Designed and implemented an end-to-end robotic control pipeline <b>integrating LLM</b> (for natural language task planning and code generation) with <b>SAM-6D</b> , enabling robotic arms to execute <b>manipulation tasks</b> from natural language commands.	
• Developed vision-language-action loop: SAM-6D processes RGB-D input for real-time object localization; LLM parses user intent and generates motion primitives; motion planner translates high-level commands to joint trajectories.	

## Selected Project

<b>Selective Memory Snapshotting for ptrace</b>	SEPT. 2025 - OCT. 2025
• Extended <b>ptrace</b> without changing its ABI by adding <b>SNAPSHOT</b> , <b>RESTORE</b> and <b>GETSNAPSHOT</b> to <b>capture/restore</b> a specified writable region of a tracee's address space.	
• Validated regions via VMA permission/length checks; stored snapshots in kernel space keyed by (pid, start, len) with caps ( <b>MAX_SNAPSHOT_LEN</b> , per-tracee <b>MAX_TOTAL_SNAPSHOT_SIZE</b> ); auto-cleanup on re-store and tracee exit.	
<b>Simulated Distributed System</b>	SEPT. 2025 - Present
• Primary-Backup One-Fault-Tolerance Storage System <ul style="list-style-type: none"><li>– Implemented with Lexical Confinement design for high-concurrency requests using Go.</li></ul>	
• Paxos-based Fault-Tolerant Key/Value Storage System <ul style="list-style-type: none"><li>– Implemented a replicated key/value store using multi-instance Paxos to totally order Get/Put/Append operations without a central coordinator.</li><li>– Ensured linearizable single-copy semantics and at-most-once execution while tolerating server and network failures, with lagging replicas catching up via the Paxos log.</li></ul>	
<b>Network File Server</b>	MAR. 2025 – APR. 2025
• Built a concurrent, crash-consistent file system with hierarchical directories, supporting <b>FS_READBLOCK</b> , <b>FS_WRITEBLOCK</b> , <b>FS_CREATE</b> , <b>FS_DELETE</b> over TCP.	

- Ensured crash safety via ordered metadata writes; scaled concurrency with Boost threads and reader-writer locks; Built network communication using POSIX sockets for client-server interaction. 

#### Memory Manager (Pager)

FEB. 2025 – MAR. 2025

- Designed a multi-process virtual memory pager with swap-backed and file-backed mappings (akin to Unix mmap), per-process page tables, and MMU protection bits.
- Handled vm\_map, page faults, fork with copy-on-write, eviction via clock (second chance), zero fill fast paths, and eager swap reservation; preserved sharing for file-backed aliases. 

#### Thread Library

JAN. 2025 – FEB. 2025

- Built a user-level thread library, managing CPU booting and thread life cycle; and monitor primitives supporting single and multiCPU execution, preemptive scheduling via timer interrupts.
- Implemented synchronization primitives like mutex, condition variable and spin-lock, using advanced Unix context management techniques. 

#### Compiler Construction - Dynamic Typed Compiler

Jan. 2025 - Apr. 2025

- Used Rust to develop a compiler for a simple language supporting dynamic typing and heap allocation on x86-64 architecture.
- Implemented front-end checking, middle-end Single Static Assignment (SSA) forming, and backend code generation with System V ABI.
- Implemented optimizations including register allocation and assertion removal.

---

### Teaching Experience

#### University of Michigan

Ann Arbor, MI

#### Foundations of Computer Science (EECS376)

FA.2025

- Knowledge includes Algorithms; Turing Reduction; P-NP problems; Cryptography
- Held Office Hours weekly.

#### Shanghai Jiao Tong University

Shanghai, China

#### Honors Calculus III (MATH2550J)

SU.2024

- Knowledge includes Linear Algebra; Multiple, Line, Surface integrals; Complex Analysis
- Held Recitation Classes and Office Hours weekly.

#### Honors Calculus II (MATH1560J)

FA. 2023

- Knowledge includes Limits and Continuity; Differentiation; Integration
- Held Recitation Classes and Office Hours weekly. (Materials Attached here 

---

### Skills

**Programming Languages:** C/C++, Python, Go, Java, Rust,etc.

**Distributed System:** Docker, Kubernetes, ZooKeeper, Resilience4j, Prometheus, ChaosBlade/ChaosMesh, etc.

**Machine Learning:** PyTorch.